Fedy BEN HASSOUNA

AI Engineer – LLM Agents, RAG & MLOps
Aix-en-Provence, France

benhassouna.fedy@gmail.com
+33 7 82 53 58 82
LinkedIn, GitHub, Portfolio

Education

National Engineering Degree – Software & Artificial Intelligence — *Tunis, Tunisia*
INSAT – National Institute of Applied Sciences and Technology — *Sept 2021 – Sept 2026*

Experience

AI Founding Engineer – *Orymate* — *Aix-en-Provence, France*
Oct 2026 – Present

- Designed the platform's **multi-agent orchestrator** (Mistral): intent routing to 4 specialized agents, confidence thresholds and rule-based fallback when the LLM is unavailable.
- Built a **multi-collection RAG** (Qdrant) across 12 marketplaces, as well as a tool-calling support agent with strict JSON output and human validation.

AI Software Engineer (Internship) – *Marketplace Distribution* — *Aix-en-Provence, France*
Mar 2026 – Sept 2026

- **Amazon Seller Central monitoring agent** (Python, CrewAI, Mistral, Selenium) across 9 accounts and 6 European markets, with financial figures kept out of the LLM: daily workload reduced from **252 to 5 min (-98%)**, **98.3% success rate** over 198 production runs.
- **AI catalog mapping tool** for Amazon, Cdiscount and Mirakl templates (200+ fields, confidence-scored mapping): a batch of 150 products goes from **∼20 h to ∼15 min**, at ∼$0.01 per product.

AI Software Engineer (Freelance) – *Oryzon Partners* — *Tunis, Tunisia*
Nov 2025 – Mar 2026

- Developed an **agentic RAG service** (streaming, authentication, caching) with an admin dashboard for document ingestion, and eBay scraping ETL pipelines with ML-based attribute extraction.

AI R&D Engineer (Internship) – *Qatar Computing Research Institute* — *Doha, Qatar*
Jul 2025 – Sept 2025

- Medical diagnosis agent validating diagnoses against patient statements: **+85% diagnostic consistency**.
- **Knowledge graph** of Statements–Diagnoses–Symptoms integrated into a multi-agent system: **-40% manual verification**.
- Led a **predictive maintenance** platform (LSTM, autoencoders, RAG agent): **-70% pump downtime**, 72% more accurate recommendations, **96% alignment** in LLM-as-judge evaluation.

Projects

Open source contribution – EleutherAI lm-evaluation-harness *(Python, LLM evaluation)* — PR #3888

- Implemented **OpenAI's SimpleQA** benchmark (4,326 questions) in the reference LLM evaluation framework, with the F1 metric matching OpenAI's official implementation; PR designated as the **reference implementation** among three competing PRs.

GAIA Agent – *LangGraph, Supabase, Gradio* — Demo

- State-machine agent solving multi-step questions from the GAIA benchmark (Hugging Face) using tools: web search, Wikipedia, arXiv and vector store.

Awards

- **3 hackathon first places**: Hack for Good 3 (Pump Watch), TN 2056 ICC TSYP12, HACK-E .
- **Selections**: MoroccoAI Summer School 2025 (50 selected out of 1,300+), MENA ML Winter School in Doha.

Skills

LLM & Agents: LangChain, LangGraph, CrewAI, Knowledge Graphs, function calling, LLM-as-judge evaluation
ML / DL: PyTorch, TensorFlow, scikit-learn, XGBoost, Hugging Face, OpenCV, YOLO, LSTM/Autoencoders
MLOps: FastAPI, Docker, GitHub Actions, Azure AI Studio, Google Cloud, Prometheus
Data: Python, SQL, PostgreSQL, MongoDB, Neo4j, Qdrant, FAISS, Pandas
Languages: Arabic (Native), French (C1), English (C1)